

Lecture Notes for Advanced Algorithms

Lecture of October 8, 2012

1 Online Algorithms and Competitive Analysis

Online algorithms are distinguished from offline algorithms by the following features:

- They do not receive all of its input at once, but instead work on partial information of requests $\sigma = \sigma(1)\sigma(2) \dots \sigma(n)$. When the online algorithm is serving the request $\sigma(i)$, it has no information about the future requests $\sigma(i')$, where $i' > i$.
- They work in a local/greedy fashion, only with the current moment and the past.
- They are not always optimal; often because it is impossible to make global assumption about the data, and the decisions that are made may turn out to not be optimal in the future.

In order to analyze online algorithms, we deploy a technique called competitive analysis, where the performance is compared with the performance of an offline optimal algorithm. The formal definition of this is:

$$C_A \leq kC_{opt} + O(1)$$

We say an algorithm is k -competitive if its cost, C_A is less or equal to the cost of the offline optimal algorithm, C_{opt} , multiplied by a positive constant k and to which is added some constant. Competitive analysis is an amortized analysis, but, instead of proving something absolute about the performance of the algorithm, it proves something relative: that the algorithm does not do worse than some constant times the performance of an optimal (usually unknown) strategy.

2 Move-To-Front : Competitive Analysis

We will run a competitive analysis of the Move-to-Front(MTF) heuristic for list update problem. In Figure 1, the oracle list is the unknown optimal list. The MTF list is




Figure 1: The general situation in the analysis of Move-To-Front.

maintained by moving the accessed element to the front of the list. As always, our potential function will quantify how poor our current structure is. In comparing our

MTF list with the oracle list on a query for x , we can partition the elements in front of x in the MTF list into two subsets: those that also occur in front of x in the oracle list, and those that occur after x in the oracle list. The former are no problem: both the oracle and the MTF heuristic have to pay one unit of actual running time for each such element; but the latter cause the access to x to run more slowly in the MTF list than in the oracle list. We will thus define the potential of x , $\phi(x)$, to be the number of elements in this second category. (Notice that we could consider the other side of the coin, namely elements that occur before x in the oracle list, but after x in the MTF list, and thus apparently yield an advantage to the MTF heuristic. We do not do it for two reasons: first, we are focused on the performance of MTF, and that is determined entirely by the elements in front of x in the MTFs list; and secondly, because it is simpler and works! Naturally, we define the potential of the MTF list as a whole as the sum of the potentials of its elements, $\Phi = \sum_x \phi(x)$.

Assume x is in the i th position in the MTF list. The real cost of accessing x is then i ; the real cost of moving x to the front of the list is constant—by keeping a pointer to the front of the list and two pointers throughout the search (to the current and the previous nodes on the list), we can move the node to the front in just three pointer changes. The change in potential due to moving x to the front is then the sum of two contributions: x will no longer contribute anything to the potential (as it will no longer have any elements in front of it), so we have a decrease of $\phi(x)$; but all $i - 1$ elements that were in front of x now have an additional element in front of them, x itself, and thus their potential might be increased by 1. This increase only occurs if x , which is now in front of them in the MTF list, is behind them in the oracle list—or, phrased in reverse, the only elements affected are those that are in front of x in the oracle list. We had partitioned these $i - 1$ elements earlier into two categories and counted $\phi(x)$ of them in the category of elements behind x in the oracle list, so there are $i - 1 - \phi(x)$ of them in front of x in the oracle list. Thus, in addition to the decrease by $\phi(x)$ due to the loss of the direct contribution of x , there is an increase by $i - 1 - \phi(x)$ due to the new indirect contribution of x , for a total change of potential of

$$-\phi(x) + (i - 1 - \phi(x)) = i - 1 - 2\phi(x)$$

Adding the real cost, we get, for a query to x in the MTF list

$$\text{real cost} + \Delta\Phi = 2 \cdot (i - \phi(x)) - 1$$

Now we must turn to the cost and change in potential in the oracle list.

2.1 Oracle: Static Version

We begin with the static version of the oracle: the list of the oracle will never be changed—it is read-only. Therefore, there is no change in the potential due the oracle itself. The real cost of the oracle is simply j , the index of the query element x in the oracle list.

To compare the oracle (j) with MTF ($2 \cdot (i - \phi(x)) - 1$), examine Figure 2. We know that there are at least $\phi(x)$ elements behind x in the oracle list; we also know that there




Figure 2: Static version of the oracle—first look.

are exactly $j - 1$ elements in front of x in the oracle list and that the list has length n . Thus we can write

$$(j - 1) + 1 + \phi(x) \leq n$$

(The one is for x itself.) Thus we can bound j by $j \leq n - \phi(x)$. But that is not the kind of bound we need: we need a lower bound for j , not an upper bound. We observed earlier that there are $i - 1 - \phi(x)$ elements that are in front of x in *both* lists: this gives us the desired lower bound: these elements are among the $j - 1$ in front of x in the oracle list, as illustrated in Figure 3, and thus we have




Figure 3: Static version of the oracle—second look.

$$j - 1 \geq i - 1 - \phi(x)$$

or, simplifying

$$j \geq i - \phi(x)$$

But now we can go back to the MTF amortized cost

$$\text{real cost} + \Delta\Phi = 2 \cdot (i - \phi(x)) - 1 \leq 2j - 1 < 2j$$

and we have shown that the amortized cost of the MTF list is strictly less than twice that of the optimal static list.

2.2 Oracle (Optimal) : Dynamic Version

Obviously, the all-knowing oracle can do better if it restructures its list between accesses. We need to look at the kind of moves it can make. First, it can move x itself,

either towards the front or towards the back. Secondly, it can move any other elements it chooses, again towards the front or towards the back. The oracle must pay for moves deeper into the list (towards the back), but moves towards the front are free—it can keep a pointer to the node in front of the place where it wants to move x . Additional real costs for the oracle go on the right-hand side of the inequality and so can only help us; potential changes go on the left-hand side, so that decreases in potential can only help, while increases in potential could damage the inequality unless they can be suitably bounded.

Let us start with moving only x . If the oracle decides to move x towards the back, say by k positions, it will incur an added cost of k , which goes on the right-hand side of the inequality we want to maintain. But it will also alter the potential, and the potential change goes on the left-hand side of the inequality, as part of our amortization. What is that change in potential? Well the potential of x itself can only decrease by such a move: there will be fewer elements behind it in the oracle list—some of the $k - 1$ elements between the original and new positions of x could be in front of x in the MTF list, as illustrated in Figure 4. The contribution $\phi(z)$ of an element z that was on front




Figure 4: Dynamic version of the oracle: x is moved towards the back.

of x in the oracle list is not affected by this move, nor is the contribution $\phi(y)$ of an element y that is behind the new position of x in the oracle list—in both cases, the sets of elements in front of, and behind z , or y , have not changed. Finally, the contribution $\phi(w)$ of an element w that is between the original and the new positions of x in the oracle list can increase by one, if x is in front of w in the MTF list. In summary, the change in potential is limited to two parts: those of the $k - 1$ elements between the old and the new positions of x in the oracle list that are in front of x in the MTF list will cause a decrease of one unit each in $\phi(x)$, and the others will see an increase of one unit each in their own ϕ contribution. There is no way for us to know how many elements are in each of these two categories, but it does not matter: the change of potential must be in the range

$$[-(k - 1), \dots, 0, \dots, k - 1]$$

and so is always less than the added real cost. Thus, if the oracle moves element x to the back, the inequality is maintained.

Let us then look at the case where the oracle moves x towards the front, illustrated in Figure 5. In this case, we are not interested in elements that are in front of x in

algorithm to an optimal strategy, despite the fact that this optimal strategy is unknown.

3 Paging Algorithms

We assume a two-level memory system, cache and main memory. When CPU executes sequence of operations $\sigma = \sigma(1)\sigma(2)\dots\sigma(n)$, it reads data from the cache. Cache memory access is very fast, but it has very limited capacity. The main memory has larger capacity, but the access of the main memory is much slower. When operation $\sigma(i)$ is asking for certain page, the processor looks for required page in the cache. If it is not found, there is a *page fault* event. Missing page has to be brought into the cache from the main memory, which is of high capacity, but slow access.

In order to maximize efficiency we need to minimize number of page faults. Many cache algorithms has been developed for this purpose. If processor needs new page, but cache is full, they decide which page should be discarded. Examples:

- FIFO (First In First Out): In this strategy, we remove the old- est page from the cache. The assumption is that old page finished its job, so it will not be used in a future. Bad side of this approach is that page will be deleted even if it is used frequently, only because of its age.
- LRU (Least Recently Used): We choose the least recently used page to be discarded. If some page was required in the close past, probably we will use it again soon.
- LFU (Least Frequently Used): We choose the page which is used the least number of times. On page fault we throw out the page with lowest counter value. Disadvantage of this strategy is that old pages are privileged because counter of the new one starts with 0.

We will illustrate the performance of the above algorithms on the following example, with cache size $k = 3$ and input set $\sigma = 23123412$.

σ	2	3	1	2	3	4	1	2
FIFO	×	×	×	✓	✓	×	✓	×
LRU	×	×	×	✓	✓	×	×	×
LFU	×	×	×	✓	✓	×	×	✓

- OPT (Optimal Offline Paging Algorithm): Discard the page for which the next request occurs the latest in a future.

σ	2	3	1	2	3	4	1	2
OPT	×	×	×	✓	✓	×	✓	✓

Theorem 1 *LFU is not c -competitive for any constant c .*

Proof In order to demonstrate that LFU is not c -competitive for any constant c , we need to find a request sequence such that, the ratio between $C_{LFU}(\sigma)$ and $C_{OPT}(\sigma)$ can

be arbitrarily large.

Let $\sigma = \overbrace{11 \dots 1}^{m \text{ times}} \overbrace{22 \dots 2}^{m \text{ times}} \dots \overbrace{kk \dots k}^{m \text{ times}} \overbrace{(k+1)(k+2)(k+1)(k+2) \dots (k+1)(k+2)}^{m \text{ times}}$, where k is the cache size. Then $C_{LFU}(\sigma) = k + 2m$ and $C_{OPT}(\sigma) = k + 2$. Since we can choose any m , there exist no constant c such that $k + 2m \leq c(k + 2) + a$. Therefore, the LFU strategy is not c -competitive for any constant c .

Theorem 2 *FIFO and LRU are k -competitive where k is the cache size.*

We would like to get an upper bound cost of FIFO and LRU and a lower bound of OPT. In order to find the desired lower and upper bounds, we recursively define *phases*. Phase 0 consists of the first k distinct pages in the request sequence. Phase $i + 1$ starts from the first page which is distinct with respect to phase i and contains the next k distinct pages. An example for $k = 3$ and the request sequence σ :

$$\sigma = \overbrace{234}^{\text{Phase 0}} \overbrace{12321}^{\text{Phase 1}} \overbrace{442443}^{\text{Phase 2}} \overbrace{1343}^{\text{Phase 3}}$$

We can verify that the following lower and upper bounds:

$$\begin{aligned} C_{FIFO}, C_{LRU} &\leq k \times \text{number of phases} \\ C_{OPT} &\geq \text{number of phases} \end{aligned}$$

In sum, we have $C_{FIFO}, C_{LRU} \leq k \times C_{OPT}$.

Competitive analyses show that the relative performances of FIFO and LRU are the same, which are better than that of LFU. But LRU achieves much better performance than FIFO in real systems. The discrepancy between theory and practice reflects that one drawback of competitive analysis is that it always assumes a malicious adversary and may ignore useful distribution information of online requests. To improve the theoretical analysis, we can design more elaborate models, or deploy average-case analysis under reasonable assumptions.