

Last Name First Name.....

Exam for Pattern Classification and Machine Learning, 2012

Prof. Matthias Seeger

- You have 180 minutes in total
- Write your name in legible characters on the top of this page
- Write all your answers on the exam sheets (no extra sheets)
- No documentation allowed apart from 1 sheet A5 of your own notes
- No calculator (or any other electronic device) is allowed
- Have your student card displayed before you on your desk
- The exam has a total of 55 points
- The exam consists of 19 pages (10 sheets, double-sided). Please check that you have received all pages, numbered 1-19. Last two empty pages can be used for scratch notes

Distribution of points:


1./7
2./4
3./6
4./5
5./6
6./7
7./6
8./4
9./10

Total:/55

1 Multilayer Perceptron (7pts)

Suppose you are given a training dataset $\{(\mathbf{x}_i, t_i) \mid i = 1, \dots, n\}$, where the input vectors $\mathbf{x}_i = [x_{i,1}, \dots, x_{i,4}]^T \in \mathbb{R}^4$, the targets $t_i \in \{-1, +1\}$. You use a multi-layer perceptron with one hidden layer of $h = 5$ units and transfer function $g(a) = \tanh(a)$, as depicted in the figure below.

Note: x_1, \dots, x_4 in the figure are components of *one* input vector $\mathbf{x} = [x_j]$.




(a) Write down the forward equations $y = y(\mathbf{x}; \mathbf{w})$ for the model (here, \mathbf{w} collects *all* parameters). This is easier if you use intermediate variables a_k and $z_k = g(a_k)$, $k = 1, \dots, 5$. (1pt)

(b) What is the total number of parameters of this model? (1pt)

(c) Suppose you have $n = 200$ datapoints and use $h = 200$ hidden units. What problem are you likely to encounter when training the MLP?

.....

How can you guard against this problem? Provide a brief explanation, and support your argument by drawing a qualitative example of the training set error and the validation set error in the figure below. **(2pts)**



(d) Your boss suggests to minimize the following error function:

$$\tilde{E} = \sum_{i=1}^n \frac{1}{5} (y(\mathbf{x}_i) - t_i)^5$$

How do you convince him that this is a bad idea? **(1pt)**

(e) You agree to use the following error function instead:

$$E = \sum_{i=1}^n E_i, \quad E_i = \frac{1}{6}(y(\mathbf{x}_i) - t_i)^6.$$

Compute the following gradient component for pattern i .

Hint: Weight $w_{54}^{(1)}$ links x_4 to a_5 . **(2pts)**


$$\frac{\partial E_i}{\partial w_{54}^{(1)}} = \dots\dots\dots$$

Space for Answer:

2 Decision Theory (4pts)

The random variables (x_1, x_2, t) , $0 \leq x_1, x_2 \leq 1$, $t \in \{0, 1\}$ are distributed as follows. x_1, x_2 are independently and uniformly drawn from the interval $[0, 1]$. Given x_1, x_2 , t is set as follows:

- If $x_1 \leq 1/3$: $\begin{cases} t = 1 & \text{if } x_2 > 1/3 \\ t = 0 & \text{if } x_2 \leq 1/3 \end{cases}$
- If $x_1 > 1/3$: $\begin{cases} t = 1 & \text{if } x_2 \leq 2/3 \\ t = 0 & \text{if } x_2 > 2/3 \end{cases}$



$(x_1, x_2) \mapsto t$ is visualized in the figure on the right.

(a) You observe x_2 only, but not x_1 . Determine the conditional probability $P(t = 1|x_2)$. What is the Bayes optimal classifier $f^*(x_2) \rightarrow \{0, 1\}$? **(2pts)**

$P(t = 1|x_2) = \dots\dots\dots$

$f^*(x_2) = \begin{cases} 0 & \text{if } \dots\dots\dots \\ 1 & \text{otherwise} \end{cases}$

Space for Derivation:

(b) In the figure above (right), shade the set $\{(x_1, x_2)\}$ where the optimal classifier f^* commits an error. What is the Bayes error of f^* ? **(2pts)**

Bayes error of f^* : $\dots\dots\dots$

3 Perceptron (6pts)

We run the perceptron algorithm in 2D, to learn a linear discriminant

$$y(x_1, x_2) = w_1x_1 + w_2x_2 + w_3, \quad \mathbf{w} = [w_1, w_2, w_3]^T.$$

The bias term is w_3 . The training dataset is $\{(\mathbf{x}_i, t_i) \mid i = 1, \dots, n\}$, where $t_i \in \{-1, +1\}$, and $\mathbf{x}_i = [x_{i1}, x_{i2}, x_{i3}]^T$, where $x_{i3} = 1$ for all i (to accommodate w_3).

(a) Complete the missing definitions (.....) in the code below. **(1pt)**

```

repeat
  for  $i \in \{1, \dots, n\}$  (in random order) do

    if ..... (HERE: Condition for update) then

      ..... (HERE: Update of weights  $\mathbf{w}$ )

    end if
  end for

until ..... (HERE: Condition for termination)

```

(b) Under which (necessary and sufficient) condition on the training set does the perceptron algorithm terminate after finitely many steps? **(1pt)**

Answer:


(c) You are given the following training dataset (also see figure on next page):

Order	x_{i1}	x_{i2}	x_{i3}	t_i
1	1	3	1	+1
2	$-\frac{3}{2}$	$-\frac{1}{2}$	1	-1
3	1	-1	1	-1
4	3	4	1	+1
5	2	$\frac{7}{2}$	1	+1


Run **three** steps of the perceptron algorithm, starting from the weight vector $\mathbf{w}_0 = [1, -1, 0]^T$, processing the datapoints $i = 1, 2, 3$ in this order.

Part (c) continues on next page


Part (c) continued



(A) After step 1



(B) After step 2



(C) After step 3

Report the weight vectors \mathbf{w}_1 , \mathbf{w}_2 , \mathbf{w}_3 after each step:

After step 1: $\mathbf{w}_1 = \dots\dots\dots$

After step 2: $\mathbf{w}_2 = \dots\dots\dots$

After step 3: $\mathbf{w}_3 = \dots\dots\dots$

Also, **draw** the corresponding separating lines into the figures (A), (B), (C) above (\times are $+1$, \circ are -1). *Hint:* The line corresponding to $\mathbf{w} = [w_1, w_2, w_3]^T$ crosses the vertical axis at $x_2 = -w_3/w_2$. **(3pts)**

Space for calculations:

(d) Describe a preprocessing technique which can speed up the convergence of the perceptron algorithm in practice. Apply the technique to the 4th datapoint $\mathbf{x}_4 = [3, 4, 1]^T$. **(1pt)**

4th point after preprocessing: $\dots\dots\dots$

Description:

4 Kernel Methods (5pts)

(a) Show that

$$K_\varepsilon(\mathbf{x}, \mathbf{y}) = [\varepsilon^2 + \mathbf{x}^T \mathbf{y}]^2, \quad \varepsilon > 0,$$

where $\mathbf{x}, \mathbf{y} \in \mathbb{R}^2$ are two-dimensional variables, is a valid kernel function that corresponds to a feature map $\phi_\varepsilon(\mathbf{x}) \in \mathbb{R}^6$.

Derive the feature mapping $\phi_\varepsilon(\mathbf{x})$. **(2pts)**

(b) A kernel method based on K_ε produces a function

$$y(\mathbf{x}) = \mathbf{w}^T \phi_\varepsilon(\mathbf{x}), \quad \mathbf{w} = [w_1, \dots, w_6]^T.$$

Show how to derive a weight vector $\tilde{\mathbf{w}} = [\tilde{w}_1, \dots, \tilde{w}_6]^T$ so that

$$\tilde{\mathbf{w}}^T \phi_1(\mathbf{x}) = y(\mathbf{x}) = \mathbf{w}^T \phi_\varepsilon(\mathbf{x}),$$

meaning that $y(\mathbf{x})$ can be represented by the kernel K_1 (with $\varepsilon = 1$) as well. **(1pt)**

(c) A friend says: “From part (b), it follows that the spaces of functions $y(\mathbf{x})$ for kernel methods using K_ε are the same for all $\varepsilon > 0$. This means that running the SVM algorithm on some data will result in the same classifier, no matter what ε is.” Explain the mistake in this argument. **(2pts)**

5 Naive Bayes Classification (6pts)

A binary Naive Bayes classifier (targets $t \in \{0, 1\}$, documents \mathbf{x}) uses seven binary features: $\phi(\mathbf{x}) \in \{0, 1\}^7$. The training data is given by

$$\Phi_0 = \begin{matrix} & \overbrace{m=1,2,3,4,5,6,7} \\ \begin{bmatrix} 1 & 0 & 1 & 1 & 1 & 0 & 1 \\ 0 & 0 & 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 0 & 1 & 1 & 0 & 1 \\ 0 & 0 & 0 & 1 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 & 1 & 1 & 1 \end{bmatrix} \end{matrix}, \quad \Phi_1 = \begin{matrix} & \overbrace{m=1,2,3,4,5,6,7} \\ \begin{bmatrix} 1 & 1 & 0 & 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 0 & 0 & 0 & 1 \\ 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 1 & 1 & 1 & 1 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 & 1 & 0 & 1 \end{bmatrix} \end{matrix}.$$

There are 7 documents for class $t = 0$, 8 documents for class $t = 1$, the feature vectors are the rows of the data matrices (Φ_0 for class 0, Φ_1 for class 1).

(a) The Naive Bayes classifier has parameters $p_m^{(k)} = \Pr\{\phi_m(\mathbf{x}) = 1 | t = k\}$. Compute the maximum likelihood estimates $\hat{p}_m^{(k)}$, $m = 1, \dots, 7$, $k = 0, 1$. **(1pt)**

m	1	2	3	4	5	6	7	
$\hat{p}_m^{(0)}$								$\hat{P}(t = 0) = \dots\dots\dots$
$\hat{p}_m^{(1)}$								$\hat{P}(t = 1) = \dots\dots\dots$

(b) A new document \mathbf{x}_* gives rise to the feature vector $\phi(\mathbf{x}_*) = [0, 0, 1, 1, 1, 0, 1]^T$. Using the trained Naive Bayes classifier, what is the probability that the document belongs to class 0? **(3pts)**

$$\hat{P}(t = 0 | \mathbf{x}_*) = \frac{1}{1 + a \cdot b^6}, \quad a = \dots\dots\dots, \quad b = \dots\dots\dots$$

Space for calculations:

(c) Your data collection collaborator informs you about a bug in the data matrices Φ_0, Φ_1 above: for features $m = 2$ and $m = 3$, all values have been flipped ($0 \leftrightarrow 1$). Determine $\hat{P}(t = 0|\mathbf{x}_*)$ for the document \mathbf{x}_* in (b) and the corrected data. **(2pts)**


$$\hat{P}(t = 0|\mathbf{x}_*) = \frac{1}{1 + c \cdot d^6}, \quad c = \dots\dots\dots, \quad d = \dots\dots\dots$$

Space for calculations:

6 Maximum Margin Perceptron and SVM (7pts)

Eleven data points representing two classes (crosses $t_i = +1$ and circles $t_i = -1$) are shown in the figure below. We use a maximum margin perceptron for classification:

$$f(\mathbf{x}) = \text{sgn}(\mathbf{w}_0^T \mathbf{x} + b), \quad \mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}.$$



(a) Draw the decision boundary of the maximum margin perceptron solution into the figure. **(1pt)**

(b) Provide the unit-norm weight vector \mathbf{w}_0 , bias parameter b and margin κ of the optimal solution. **(3pts)**

$\mathbf{w}_0 = \dots\dots\dots$

$b = \dots\dots\dots$

$\kappa = \dots\dots\dots$

Space for calculations

(c) How many support vectors are there in the optimal solution? Provide their indices. **(1pt)**

.....

(d) Express the normalized weight vector \mathbf{w}_0 of (b) as linear combination of the support vectors found in (c). **(2pts)**


Hint: Consider the orthogonal projection of the class -1 support vector onto the line through the class $+1$ support vectors.

Result: $\mathbf{w}_0 =$

Space for calculations

7 Least Squares Regression (6pts)

You are given the following dataset $\{(x_1, t_1), \dots, (x_5, t_5)\}$, where $x_i, t_i \in \mathbb{R}$.



You choose a linear model

$$y(x) = ax + b.$$

To fit the model, you find $a, b \in \mathbb{R}$ which minimize the squared error function:

$$E(a, b) = \frac{1}{2} \sum_{i=1}^n (y(x_i) - t_i)^2.$$

(a) Begin by computing the following statistics from the data: **(1pt)**

$$\langle x \rangle = \frac{1}{n} \sum_{i=1}^n x_i = \dots\dots\dots$$

$$\langle x^2 \rangle = \frac{1}{n} \sum_{i=1}^n x_i^2 = \dots\dots\dots$$

$$\langle xt \rangle = \frac{1}{n} \sum_{i=1}^n x_i t_i = \dots\dots\dots$$

$$\langle t \rangle = \frac{1}{n} \sum_{i=1}^n t_i = \dots\dots\dots$$

Space for calculations:

(b) Derive the gradient components $\frac{\partial E}{\partial a}$ and $\frac{\partial E}{\partial b}$. Express them as functions of $\langle x \rangle$, $\langle x^2 \rangle$, $\langle xt \rangle$ and $\langle t \rangle$. **(2pts)**

$$\frac{\partial E}{\partial a} = \dots\dots\dots$$

$$\frac{\partial E}{\partial b} = \dots\dots\dots$$

Derivation (more space on following page):

Derivation for (b), continued:

(c) Compute the global minimum point (a_*, b_*) of the error $E(a, b)$. Draw the line corresponding to your solution into the figure on the previous page. **(2pts)**

$a_* = \dots\dots\dots$ $b_* = \dots\dots\dots$

Calculation:

(d) Suppose you choose the model $y(x) = cx^2 + ax + b$ instead, with parameters $a, b, c \in \mathbb{R}$. Mark the correct answer from the following options, and **provide a brief explanation**. **(1pt)**


“The minimum squared error for the new setup, compared to the minimum error worked out above, ...”

- ☐ stays the same
- ☐ increases or stays the same
- ☐ decreases or stays the same
- ☐ can decrease or increase, depending on the data

Explanation:

8 Principal Components Analysis (4pts)

(a) How do you determine the *first* (leading) principal components direction \mathbf{u} for data $\mathbf{x}_1, \dots, \mathbf{x}_n$, where $\mathbf{x}_i \in \mathbb{R}^p$? Provide the definition of the matrix you use. (1pt)



(b) Consider the six datapoints in \mathbb{R}^2 , depicted in the figure above. What is the *first* (leading) principal components direction \mathbf{u} ? What is the corresponding eigenvalue λ ? *Hint:* Consider the orthogonal directions $(1/\sqrt{2})[1, -1]^T$ and $(1/\sqrt{2})[1, 1]^T$. (3pts)

$\mathbf{u} = \dots\dots\dots$

$\lambda = \dots\dots\dots$

Space for calculations on the following page

Space for calculations

9 Maximum Likelihood (10pts)

We model n positive data points $x_i > 0$ ($1 \leq i \leq n$) as drawn independently from a probability distribution with density function

$$p(x|\gamma) = \frac{1}{2}\gamma^3 x^2 \exp(-\gamma x) \quad \text{for } x > 0 \quad (1)$$

and $p(x|\gamma) = 0$ for $x \leq 0$. Here, $\gamma > 0$ is a parameter.

(a) Write down the likelihood function for the model (1) and the data. **(1pt)**

(b) Find the optimal value $\hat{\gamma}$ of γ by the principle of maximum likelihood. Write the result in the form **(2pts)**

$$\frac{1}{\hat{\gamma}} = \dots\dots\dots$$

Derivation:

(c) Write down a mixture model with four components $p(x|\omega_k)$, each of the form of Eq. 1, but with different γ_k . Denote the prior probabilities for the different components by $P(\omega_k)$, $k = 1, \dots, 4$. **(1pt)**

$$p_{\text{mixture}}(x) = \dots\dots\dots$$

(d) How many free and independent parameters does your mixture model of part (c) have? **(1pt)**

(e) For the mixture model in part (c), we have that $P(\omega_k) = 1/4$, $k = 1, \dots, 4$, $\gamma_1 = 2$, and $\gamma_2 = \gamma_3 = \gamma_4 = 1$. Compute the posterior probability $P(\omega_1|x_i)$ of x_i coming from component 1, for $x_i = \log 2$ (natural logarithm: $e^{x_i} = 2$). **(2pts)**

$P(\omega_1|x_i = \log 2) =$

Space for calculations:

(f) We want to estimate parameters by maximum likelihood. Derive an update equation for the parameter γ_1 of mixture component 1. To do so, define

$P(\omega_k|x_i) = \frac{p(x_i|\omega_k)P(\omega_k)}{p(x_i)}$ and show that for a constant C (provide it!) **(3pts)**

$$\frac{1}{\hat{\gamma}_1} = C \frac{\sum_i P(\omega_1|x_i)x_i}{\sum_i P(\omega_1|x_i)}$$

Note: Just determine the stationary point (no need for 2nd derivative).

Derivation:

Additional space for notes:

Additional space for notes: